

Prometheus - Query Samples

Eg: `<metric-name>{<label-name> != "/"}/ 1e9 > 100`

Querying

All partitions in my entire infrastructure with more than 100GB capacity that are not mounted on root?

```
node_filesystem_bytes_total{mountpoint!="/"}/ 1e9 > 100
```

{device="sda1", mountpoint="/home", instance="10.0.0.1"}	118.8
{device="sda1", mountpoint="/home", instance="10.0.0.2"}	118.8
{device="sdb1", mountpoint="/data", instance="10.0.0.2"}	451.2
{device="xdrv", mountpoint="/mnt", instance="10.0.0.3"}	320.0

Sum By all paths

What's the ratio of request errors across all service instances?

```
sum(rate(http_requests_total{status="500"}[5m]))  
/ sum(rate(http_requests_total[5m]))
```

```
{} 0.029
```

Group by paths

What's the ratio of request errors across all service instances?

```
sum by(path) (rate(http_requests_total{status="500"}[5m]))  
/ sum by(path) (rate(http_requests_total[5m]))
```

{path="/status"}	0.0039
{path="/"}	0.0011
{path="/api/v1/topics/:topic"}	0.087
{path="/api/v1/topics"}	0.0342

Latency Distribution

99th percentile request latency across all instances?

```
histogram_quantile(0.99,  
  sum without(instance) (rate(request_latency_seconds_bucket[5m]))  
)
```

{path="/status", method="GET"}	0.012
{path="/", method="GET"}	0.43
{path="/api/v1/topics/:topic", method="POST"}	1.31
{path="/api/v1/topics", method="GET"}	0.192

Prometheus - Query Samples

Eg: `<metric-name>{<label-name> != "/" } / 1e9 > 100`

Querying

All partitions in my entire infrastructure with more than 100GB capacity that are not mounted on root?

```
node_filesystem_bytes_total{mountpoint!="/" } / 1e9 > 100
```

{device="sda1", mountpoint="/home", instance="10.0.0.1"}	118.8
{device="sda1", mountpoint="/home", instance="10.0.0.2"}	118.8
{device="sdb1", mountpoint="/data", instance="10.0.0.2"}	451.2
{device="xdrv", mountpoint="/mnt", instance="10.0.0.3"}	320.0

Sum By all paths

What's the ratio of request errors across all service instances?

```
sum(rate(http_requests_total{status="500"}[5m]))  
/ sum(rate(http_requests_total[5m]))
```

{}	0.029
----	-------

Group by paths

What's the ratio of request errors across all service instances?

```
sum by(path) (rate(http_requests_total{status="500"}[5m]))  
/ sum by(path) (rate(http_requests_total[5m]))
```

{path="/status"}	0.0039
{path="/"}	0.0011
{path="/api/v1/topics/:topic"}	0.087
{path="/api/v1/topics"}	0.0342

Latency Distribution

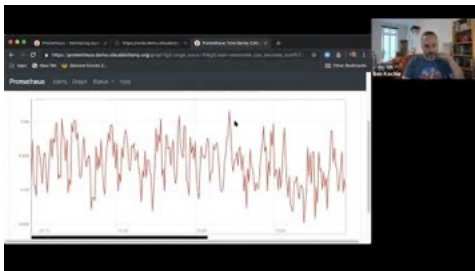
99th percentile request latency across all instances?

```
histogram_quantile(0.99,  
  sum without(instance) (rate(request_latency_seconds_bucket[5m]))  
)
```

{path="/status", method="GET"}	0.012
{path="/", method="GET"}	0.43
{path="/api/v1/topics/:topic", method="POST"}	1.31
{path="/api/v1/topics", method="GET"}	0.192

Prometheus - Reference Links

[Prometheus 101](#)



[Setting up Prometheus and Grafana for monitoring your servers](#)

